

02-Models-Performance

March 1, 2021

1 Model's Performance

We need to evaluate of our model's performance after training procedure. In machine learning, a common task is the study and construction of algorithms that can learn from and make predictions on data. The data used to build final model usually comes from multiple datasets. In particular, three datasets are commonly used in different stages of the creation of the model.

We can perform this task by dividing our dataset into 3 parts.

1.0.1 Training Set

The model is initially fit on a training set. That is a set of examples used to fit the parameters of the model. We can denote the training dataset as

$$(X, y) = (x^1, y^1), (x^2, y^2), \dots, (x^m, y^m)$$

1.0.2 Validation Set

The validation dataset provides an unbiased evaluation of a model fit on the training dataset while tuning the model's hyperparameters. For example, a set of examples used to tune the parameters of a classifier, in the MLP case, we would use the validation set to find the “optimal” number of hidden units or determine a stopping point for the back-propagation algorithm

1.0.3 Test Set

The test set dataset is a dataset used to provide an unbiased evaluation of a final model fit on the training dataset. We can compute errors from this dataset.

Let's show training/validation/test split.

```
[5]: import pandas as pd
df = pd.read_csv('./data/Credit.csv')
df.drop(columns=['Unnamed: 0'], inplace=True)
df.head()
```

```
[5]:
```

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	\
0	14.891	3606	283	2	34	11	Male	No	Yes	
1	106.025	6645	483	3	82	15	Female	Yes	Yes	
2	104.593	7075	514	4	71	11	Male	No	No	
3	148.924	9504	681	3	36	11	Female	No	No	
4	55.882	4897	357	2	68	16	Male	No	Yes	

	Ethnicity	Balance
0	Caucasian	333
1	Asian	903
2	Asian	580
3	Asian	964
4	Caucasian	331

```
[6]: print('Number of the samples: {0}'.format(len(df)))
```

Number of the samples: 400

```
[7]: def trainTestSplit(data, ratio = 0.8):
    if isinstance(data, pd.DataFrame):
        data = data.sample(frac=1).reset_index(drop=True)
        train_pct_index = int(ratio * len(data))
        train = data.iloc[:train_pct_index,:]
        test = data.iloc[train_pct_index:,:]
        test.reset_index(inplace=True, drop = True)
        return train, test
    elif isinstance(data,np.array):
        X_train, X_test = data[:train_pct_index,0], data[train_pct_index:,0]
        Y_train, Y_test = data[:train_pct_index,1:], data[train_pct_index:,1:]
        return X_train, X_test, Y_train, Y_test
```

```
[8]: train, test = trainTestSplit(df,0.6)
```

```
[9]: train.head()
```

```
[9]:
```

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	\
0	53.598	3714	286	3	73	17	Female	No	Yes	
1	26.427	5533	433	5	50	15	Female	Yes	Yes	
2	30.111	4336	339	1	81	18	Male	No	Yes	
3	39.609	2539	188	1	40	14	Male	No	Yes	
4	24.543	3206	243	2	62	12	Female	No	Yes	

	Ethnicity	Balance
0	African American	0
1	Asian	1404
2	Caucasian	347
3	Asian	0

4 Caucasian 95

```
[10]: test.head()
```

```
[10]:
```

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	\
0	59.879	6906	527	6	78	15	Female	No	No	
1	39.055	5565	410	4	48	18	Female	No	Yes	
2	36.362	5183	376	3	49	15	Male	No	Yes	
3	49.927	6396	485	3	75	17	Female	No	Yes	
4	62.413	6457	455	2	71	11	Female	No	Yes	

	Ethnicity	Balance
0	Caucasian	1032
1	Caucasian	772
2	African American	654
3	Caucasian	890
4	Caucasian	762

```
[11]: print('Number of samples in training set: {0} \nNumber of samples in test set: \n\n{1}'.format(len(train),len(test)))
```

Number of samples in training set: 240

Number of samples in test set: 160

Now let's split the test set as validation/test.

```
[12]: test, val = trainTestSplit(test,0.5)
```

```
[13]: test.head()
```

```
[13]:
```

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	\
0	26.400	5640	398	3	58	15	Female	No	No	
1	12.000	4160	320	4	28	14	Female	No	Yes	
2	16.482	3326	268	4	41	15	Male	No	No	
3	16.711	5274	387	3	42	16	Female	No	Yes	
4	29.638	5833	433	3	29	15	Female	No	Yes	

	Ethnicity	Balance
0	Asian	905
1	Caucasian	602
2	Caucasian	271
3	Asian	863
4	Asian	942

```
[14]: val.head()
```

```
[14]:
```

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	\
0	10.503	2923	232	3	25	18	Female	No	Yes	

1	115.123	7760	538	3	83	14	Female	No	No
2	128.669	9824	685	3	67	16	Male	No	Yes
3	39.055	5565	410	4	48	18	Female	No	Yes
4	61.620	5140	374	1	71	9	Male	No	Yes

	Ethnicity	Balance
0	African American	191
1	African American	661
2	Asian	1243
3	Caucasian	772
4	Caucasian	302

```
[15]: print('Number of samples in test set: {0} \nNumber of samples in validation set:
      ↪ {1}'.format(len(test),len(val)))
```

Number of samples in test set: 80

Number of samples in validation set: 80

1.1 Evaluating Regression Models

1.1.1 Evaluating Linear Regression

After we built a Linear Regression model, we need to see evaluation metrics based on this model due to validity of our model.

The L^2 loss, mean squared error (MSE), is one of the tools for evaluation. We defined last week as:

$$\begin{aligned}
 L(\hat{y}, y) &= \frac{1}{m} \sum_i [y^i - (\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n)]^2 \\
 &= \frac{1}{m} \sum_i (y^i - \hat{y}^i)^2
 \end{aligned}$$

MSE or Mean Squared Error is one of the most preferred metrics for regression tasks. It is simply the average of the squared difference between the target value and the value predicted by the regression model. As it squares the differences, it penalizes even a small error which leads to over-estimation of how bad the model is.

If MSE is relatively big, that mean our model is not suitable for the data or vice-versa.

Also the Root Mean Squared Error (RMSE) can be used for evaluation. RMSE is the most widely used metric for regression tasks and is the square root of the averaged squared difference between the target value and the value predicted by the model. It is preferred more in some cases because the errors are first squared before averaging which poses a high penalty on large errors. This implies that RMSE is useful when large errors are undesired.

$$L(\hat{y}, y) = \sqrt{\frac{\sum_i (y^i - \hat{y}^i)^2}{m}}$$

Another example for evaluation is Mean Absolute Error (MAE). MAE is the absolute difference between the target value and the value predicted by the model. The MAE is more robust to outliers and does not penalize the errors as extremely as mse. MAE is a linear score which means all the individual differences are weighted equally. It is not suitable for applications where you want to pay more attention to the outliers.

$$L(\hat{y}, y) = \frac{1}{m} \sum_i |y^i - \hat{y}^i|$$

Another powerfull example for evaluation is R^2 error (also known as the Coefficient of Determination). The MSE provides an absolute measure of the lack of fit of the model to the data. But since it is measured in the units of y , it is not always clear what constitutes a good MSE. The R^2 statistics provides an alternative measure of fit. It takes the form of proportion and so it always takes on a value between 0 and 1, and it is independent of the scale of Y .

To calculate R^2 , we use the formula

$$\begin{aligned} R^2 &= \frac{\sum_i (y^i - \bar{y})^2 - \sum_i (y^i - \hat{y}^i)^2}{\sum_i (y^i - \bar{y})^2} \\ &= 1 - \frac{\sum_i (y^i - \hat{y}^i)^2}{\sum_i (y^i - \bar{y})^2} \end{aligned}$$

where,

$$\bar{y} = \frac{1}{m} \sum_i y^i$$

An R^2 statistic, if feature X can predict the target, then the proportion is high and the R^2 value will be close to 1. If opposite is true, the R^2 value is then closer to 0.

```
[1]: import numpy as np # array (dizi)
import os # file system
import matplotlib.pyplot as plt # data visualization
import pandas as pd # dataframe

# scikitlearn
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

pd.options.mode.chained_assignment = None

df = pd.read_csv('./data/slr.csv')
df.head()
```

```
# PCA -> principal component analysis
```

```
[1]:      SAT   GPA
0  1714  2.40
1  1664  2.52
2  1760  2.54
3  1685  2.74
4  1693  2.83
```

```
[2]: df.describe()
```

```
[2]:          SAT      GPA
count    84.000000  84.000000
mean    1845.273810   3.330238
std      104.530661   0.271617
min     1634.000000   2.400000
25%     1772.000000   3.190000
50%     1846.000000   3.380000
75%     1934.000000   3.502500
max     2050.000000   3.810000
```

```
[3]: def trainTestSplit(data, ratio = 0.8):
      if isinstance(data, pd.DataFrame):
          data = data.sample(frac=1).reset_index(drop=True)
          train_pct_index = int(ratio * len(data))
          train = data.iloc[:train_pct_index,:]
          test = data.iloc[train_pct_index:,:]
          test.reset_index(inplace=True, drop = True)
          return train, test
      elif isinstance(data,np.array):
          X_train, X_test = data[:train_pct_index,0], data[train_pct_index:,0]
          Y_train, Y_test = data[:train_pct_index,1:], data[train_pct_index:,1:]
          return X_train, X_test, Y_train, Y_test
```

```
[4]: df_train, df_test = trainTestSplit(df, ratio = 0.9)
```

```
[5]: print('Number of samples in data: {0} \nNumber of samples in training set: {1}\n
      ↪\nNumber of samples in test set: {2}'
      .format(len(df),len(df_train),len(df_test)))
```

```
Number of samples in data: 84
Number of samples in training set: 75
Number of samples in test set: 9
```

```
[6]: X_train = df_train.iloc[:,0]
      y_train = df_train.iloc[:,1]
```

```
[7]: X_train_arr = np.array(X_train) #  
X_train_arr.shape = (X_train_arr.shape[0],1)  
y_train_arr = np.array(y_train)  
y_train_arr.shape = (y_train_arr.shape[0],1)
```

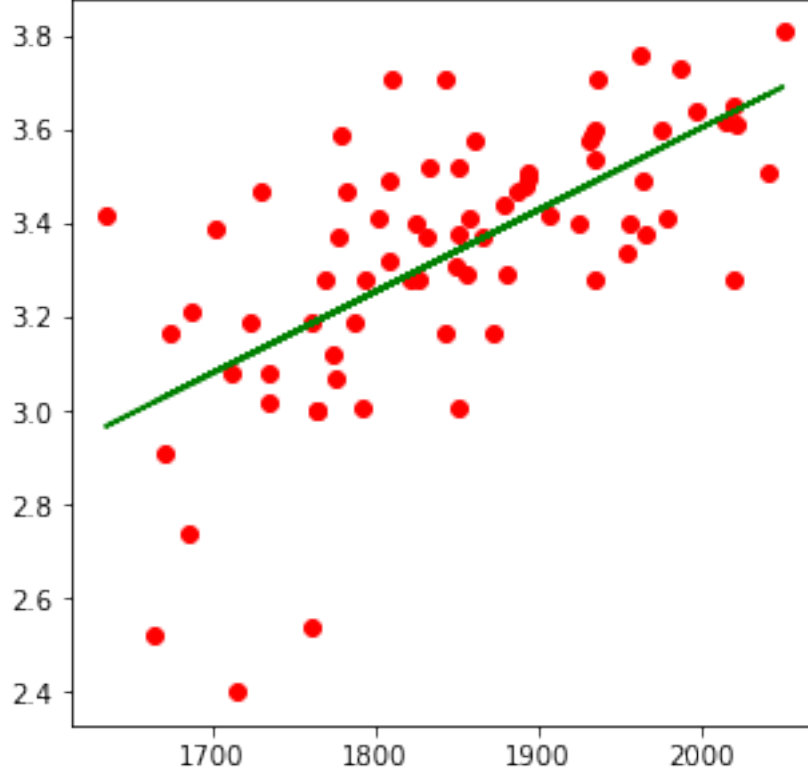
```
model = LinearRegression();  
model.fit(X_train_arr, y_train_arr)  
  
y_pred = model.predict(X_train_arr);
```

```
[10]: from sklearn.metrics import r2_score  
mse = mean_squared_error(y_train_arr,y_pred)  
r2 = r2_score(y_train_arr,y_pred)  
print(mse,r2)
```

```
0.045316419096650816 0.4120723365154918
```

```
[11]: import matplotlib.pyplot as plt  
fig = plt.figure(figsize=(5,5))  
plt.scatter(X_train, y_train, color = "red")  
plt.plot(X_train, y_pred, color = "green")
```

```
[11]: [<matplotlib.lines.Line2D at 0x7f5884c0dac8>]
```



1.1.2 Evaluating Logistic Regression

Misclassification Error Before evaluating Logistic Regression, let's do a probabilistic evaluation of errors on decision making.

Suppose that we have two classes $C = \{C_k : k \in \{1, 2\}\}$ and we have random samples from a Gaussian Distribution, x . These samples are generated from two Gaussian Distributions. Let's denote these two Gaussian Distributions as $C_1 \sim \mathcal{N}(\mu_1, \sigma_1)$ and $C_2 \sim \mathcal{N}(\mu_2, \sigma_2)$. Some samples are generated from the first distribution that correspond to class C_1 and others are generated from the second distribution that correspond to class C_2 . As you can see, that implies, this is a binary classification task. In real life, our generated samples x are not real data. But it is clear to see the concepts.

Now let's define our distributions. First, let's choose the distribution of C_1 to be bimodal. In other words, it is the concatenation of two Gaussian Distributions. Choose parameters for C_1 , $\mu_{11} = -2$, $\mu_{12} = 25$, $\sigma_{11} = 5$, $\sigma_{12} = 7$. So, the distribution of C_1 is,

$$p(x, C_{1_1}) = \frac{1}{5\sqrt{2\pi}} \exp\left(\frac{-(x - (-2))^2}{2 \cdot 25}\right)$$

$$p(x, C_{1_2}) = \frac{1}{7\sqrt{2\pi}} \exp\left(\frac{-(x - 25)^2}{2 \cdot 49}\right)$$

$$p(x, C_1) = \left[\frac{1}{5\sqrt{2\pi}} \exp\left(\frac{-(x - (-2))^2}{2 \cdot 25}\right); \frac{1}{7\sqrt{2\pi}} \exp\left(\frac{-(x - 25)^2}{2 \cdot 49}\right) \right]$$

And the distribution of C_2 as follows,

$$p(x, C_2) = \frac{1}{6\sqrt{2\pi}} \exp\left(\frac{-(x - 25)^2}{2 \cdot 36}\right)$$

```
[29]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

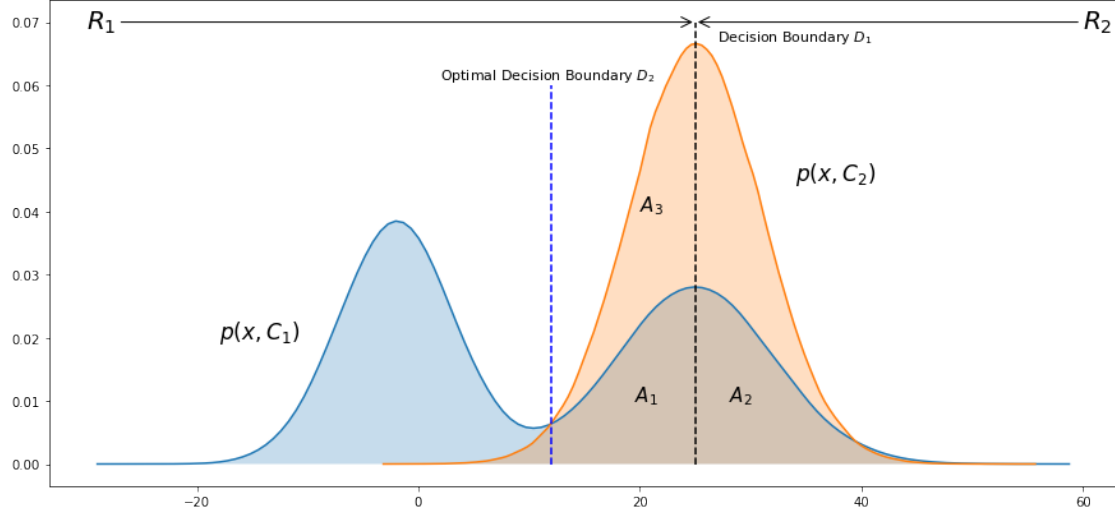
mean2, std2 = 25, 6

mean11, std11 = -2, 5
mean12, std12 = 25, 7
X1 = np.random.normal(mean11, std11, 100000)
X2 = np.random.normal(mean12, std12, 100000)
dist1 = np.concatenate([X1, X2])

dist2 = np.random.normal(mean2, std2, 750000)

fig = plt.figure(figsize = (15,7))
sns.kdeplot(dist1, shade=True)
sns.kdeplot(dist2, shade=True)
x1, y1 = [25, 25], [0, 0.07]
x2, y2 = [12,12], [0, 0.06]
plt.plot(x2,y2,linestyle='dashed',color='blue')
plt.plot(x1,y1,linestyle='dashed',color='black')
plt.text(19.5, 0.01, '$A_1$', fontsize=15)
plt.text(28, 0.01, '$A_2$', fontsize=15)
plt.text(20, 0.04, '$A_3$', fontsize=15)
plt.annotate('$R_1$', xy=(25,0.07), xytext=(-30,0.07),
            arrowprops={'arrowstyle': '->'}, va='center', fontsize=20)
plt.annotate('$R_2$', xy=(25,0.07), xytext=(60,0.07),
            arrowprops={'arrowstyle': '->'}, va='center', fontsize=20)
plt.text(27, 0.067, 'Decision Boundary $D_1$', fontsize=11)
plt.text(2, 0.061, 'Optimal Decision Boundary $D_2$', fontsize=11)
plt.text(-18, 0.02, '$p(x, C_1)$', fontsize=17)
plt.text(34, 0.045, '$p(x, C_2)$', fontsize=17)
plt.plot()
```

[29]: []



We need a rule that assigns each value of x to one of the available classes. Such a rule will divide the input space into regions R_k called *decision regions*, one for each class, such that all points in R_k are assigned into class C_k . The boundaries between decision regions are called *decision boundaries*.

A mistake occurs when an input vector belonging to class C_1 is assigned to class C_2 or vice versa. The probability of this occurrence is given by:

$$p(\text{mistake}) = p(x \in R_1, C_2) + p(x \in R_2, C_1)$$

$$= \int_{R_1} p(x, C_2) dx + \int_{R_2} p(x, C_1) dx$$

Clearly to minimize $p(\text{mistake})$, we should arrange that each x is assigned to whichever class has smaller value of in the integrals. Thus, if $p(x, C_1) > p(x, C_2)$ for a given value of x , then we should assign that x to class C_1 . From the product rule of probability, we have $p(x, C_k) = p(C_k|x)p(x)$. As you can see, both sides are divided by factor $p(x)$ which is common. So we can rewrite the inequality as $p(C_1|x) > p(C_2|x)$. This means that the minimum probability of making a mistake is obtained if each value of x is assigned to the class for which the posterior probability $p(C_k|x)$ is largest.

For the general case of multi-class tasks, we can easily say that maximizing the correctness is slightly easier than minimizing the mistake

$$p(\text{correct}) = \sum_{k=1}^K p(x \in R_k, C_k)$$

$$= \sum_{k=1}^K \int_{R_k} p(x, C_k) dx$$

To explain the above figure, values of $x > D_1$ are classified as class C_2 and hence belong to decision region R_2 . Whereas points $x < D_1$ classified as C_1 and belong to R_1 . The mistake/error arise from regions A_1, A_2 and A_3 . So that; for $x < D_1$, the errors are due to points from class C_2 being misclassified as C_1 (computed as sum $A_1 + A_3$). And; for $x > D_1$, the errors are due to points from class C_1 being misclassified as C_2 (area of A_2).

Actually, as you can see, the optimal choice for decision boundary is line D_2 . Because in this case, region A_3 disappears. This is equivalent to the minimum misclassification rate decision rule, which assigns each value of x to the class having the higher posterior probability $p(C_k|x)$.

Now let's apply what we learn above to evaluating Logistic regression.

Confusion Matrix As in the Linear Regression, the binary cross entropy loss can be used for model evaluation.

$$L(\hat{y}, y) = - \sum_i^m y^i \log(p^i) + (1 - y^i) \log(1 - p^i) = - \sum_i^m y^i \log(\hat{y}^i) + (1 - y^i) \log(1 - \hat{y}^i)$$

But as in linear regression, loss is a relative metric. It cannot say anything about our model's validity or *accuracy*.

The confusion matrix is used for this evaluation:

In the field of machine learning and specifically the problem of statistical classification, a confusion matrix, also known as an error matrix, is a specific table layout that allows visualization of the performance of an algorithm, typically a supervised learning one (in unsupervised learning it is usually called a matching matrix). Each row of the matrix represents the instances in a predicted class while each column represents the instances in an actual class (or vice versa). The name stems from the fact that it makes it easy to see if the system is confusing two classes (i.e. commonly mislabeling one as another).

```
[2]: from IPython.display import Image
      from IPython.core.display import HTML
      Image(filename= "./img/cm.png",width=500, height=500)
```

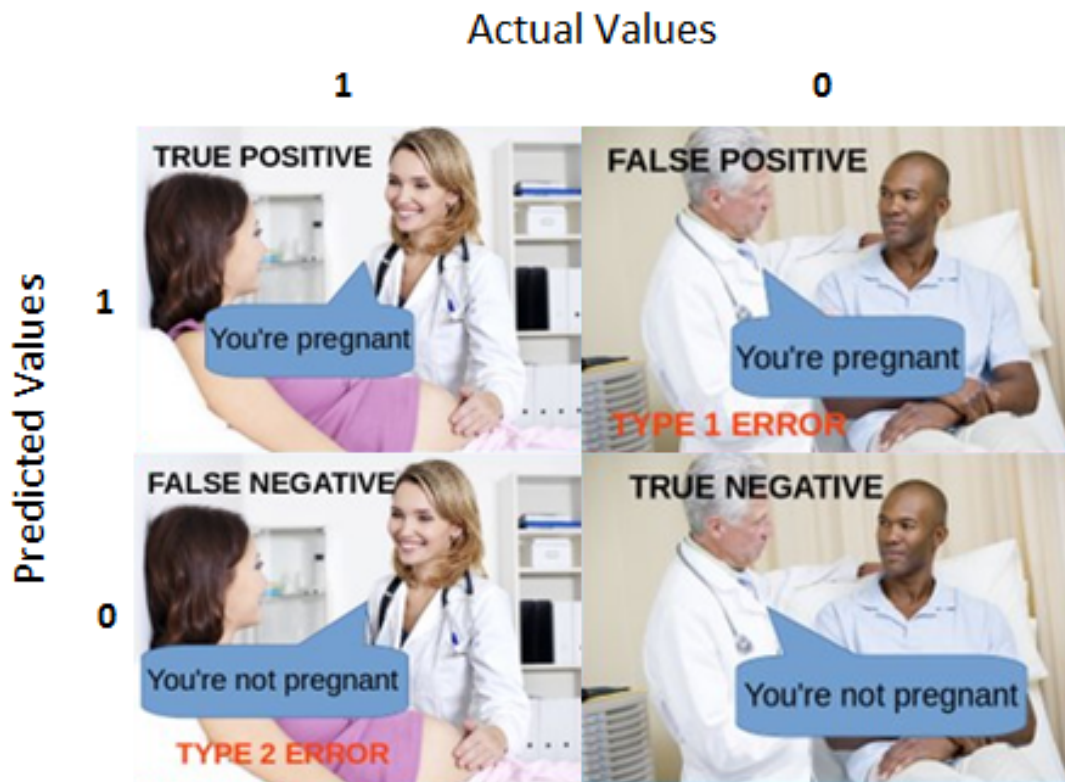
[2]:

		Predicted class	
		+	−
Actual class	+	TP True Positives	FN False Negatives Type II error
	−	FP False Positives Type I error	TN True Negatives

- **True Positive (TP):** For a given classes $C = \{C_0 = 0, C_1 = 1\}$, the input is predicted as $C_1 = 1$ and actual class of input is $C_1 = 1$.
- **True Negative (TN):** For a given classes $C = \{C_0 = 0, C_1 = 1\}$, the input is predicted as $C_0 = 0$ and actual class of input is $C_0 = 0$.
- **False Positive (FP, Type 1 Error):** For a given classes $C = \{C_0 = 0, C_1 = 1\}$, the input is predicted as $C_1 = 1$ but actual class of input is $C_0 = 0$.
- **False Negative (FN, Type 2 Error):** For a given classes $C = \{C_0 = 0, C_1 = 1\}$, the input is predicted as $C_0 = 0$ but actual class of input is $C_1 = 1$.

```
[3]: from IPython.display import Image
from IPython.core.display import HTML
Image(filename= "./img/cm_2.png",width=500, height=500)
```

[3]:



But TP, TN, FP, FN are not telling us informations about model's performance individually. Let's introduce metrics based on this values.

- **Accuracy:**

$$\frac{TP + TN}{TP + TN + FP + FN}$$

is overall performance of model. Ratio between true predictions and true predictions plus false predictions.

- **Precision:**

$$\frac{TP}{TP + FP}$$

it measures how accurate the positive predictions are.

Precision talks about how precise/accurate your model is out of those predicted positive, how many of them are actual positive. Precision is a good measure to determine, when the costs of False Positive is high. For instance, email spam detection. In email spam detection, a false positive means that an email that is non-spam (actual negative) has been identified as spam (predicted spam). The email user might lose important emails if the precision is not high for the spam detection model.

- **Recall:**

$$\frac{TP}{TP + FN}$$

it measures out of all the positive classes, how much we predicted correctly. It should be high as possible. Coverage of actual positive sample

Recall actually calculates how many of the Actual Positives our model capture through labeling it as Positive (True Positive). Applying the same understanding, we know that Recall shall be the model metric we use to select our best model when there is a high cost associated with False Negative.

For instance, in fraud detection or sick patient detection. If a fraudulent transaction (Actual Positive) is predicted as non-fraudulent (Predicted Negative), the consequence can be very bad for the bank.

- **F1 score:**

$$2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = 2 \times \frac{\frac{TP^2}{TP^2 + TP \times FP + FN \times TP + FN \times FP}}{\frac{TP^2 + FN \times TP + TP^2 + FP \times TP}{TP^2 + TP \times FP + FN \times TP + FN \times FP}} = 2 \times \frac{TP^2}{2TP^2 + FN \times TP + FP \times TP} = \frac{2TP^2}{2TP^2 + FN \times TP + FP \times TP}$$

It is difficult to compare two models with low precision and high recall or vice versa. So to make them comparable, we use F-Score. F-score helps to measure Recall and Precision at the same time. It uses Harmonic Mean in place of Arithmetic Mean by punishing the extreme values more. F1 Score might be a better measure to use if we need to seek a balance between Precision and Recall AND there is an uneven class distribution (large number of Actual Negatives).

Let's do an example.

```
[6]: from IPython.display import Image
     from IPython.core.display import HTML
     Image(filename= "./img/cm_example.jpeg",width=500, height=500)
```

[6]:

```
[12]: from sklearn.linear_model import LogisticRegression
import pandas as pd
import numpy as np
import warnings

warnings.filterwarnings("ignore", category=DeprecationWarning)

def trainTestSplit(data, ratio = 0.8):
    if isinstance(data, pd.DataFrame):
        data = data.sample(frac=1).reset_index(drop=True)
        train_pct_index = int(ratio * len(data))
        train = data.iloc[:train_pct_index,:]
        test = data.iloc[train_pct_index:,:]
        test.reset_index(inplace=True, drop = True)
        return train, test
    elif isinstance(data,np.array):
        X_train, X_test = data[:train_pct_index,0], data[train_pct_index:,0]
        Y_train, Y_test = data[:train_pct_index,1:], data[train_pct_index:,1:]
        return X_train, X_test, Y_train, Y_test

df = pd.read_csv('./data/diabetes.csv')
df.head()
```

```
[12]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	2	138	62	35	0	33.6	
1	0	84	82	31	125	38.2	
2	0	145	0	0	0	44.2	
3	0	135	68	42	250	42.3	
4	1	139	62	41	480	40.7	

	DiabetesPedigreeFunction	Age	Outcome
0	0.127	47	1
1	0.233	23	0
2	0.630	31	1
3	0.365	24	1
4	0.536	21	0

```
[13]: df.describe()
```

```
[13]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	3.703500	121.182500	69.145500	20.935000	80.254000	
std	3.306063	32.068636	19.188315	16.103243	111.180534	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	63.500000	0.000000	0.000000	
50%	3.000000	117.000000	72.000000	23.000000	40.000000	

75%	6.000000	141.000000	80.000000	32.000000	130.000000
max	17.000000	199.000000	122.000000	110.000000	744.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	2000.000000	2000.000000	2000.000000	2000.000000
mean	32.193000	0.470930	33.090500	0.342000
std	8.149901	0.323553	11.786423	0.474498
min	0.000000	0.078000	21.000000	0.000000
25%	27.375000	0.244000	24.000000	0.000000
50%	32.300000	0.376000	29.000000	0.000000
75%	36.800000	0.624000	40.000000	1.000000
max	80.600000	2.420000	81.000000	1.000000

```
[14]: train, test = trainTestSplit(df,0.8)
print('Number of samples in data: {0} \nNumber of samples in training set: {1}\nNumber of samples in test test: {2}'
      .format(len(df),len(train),len(test)))
```

Number of samples in data: 2000
 Number of samples in training set: 1600
 Number of samples in test test: 400

```
[15]: X_train = np.array(train.iloc[:, :-1])
y_train = np.array(train.iloc[:, -1])
X_test = np.array(test.iloc[:, :-1])
y_test = np.array(test.iloc[:, -1])

with warnings.catch_warnings():
    warnings.simplefilter("ignore")
    model = LogisticRegression();
    model.fit(X_train, y_train);
```

```
[16]: score_train = model.score(X_train, y_train)
score_test = model.score(X_test, y_test)
print('Accuracy Train: {0} \nAccuracy Test: {1}'.format(score_train,score_test))
```

Accuracy Train: 0.78375
 Accuracy Test: 0.76

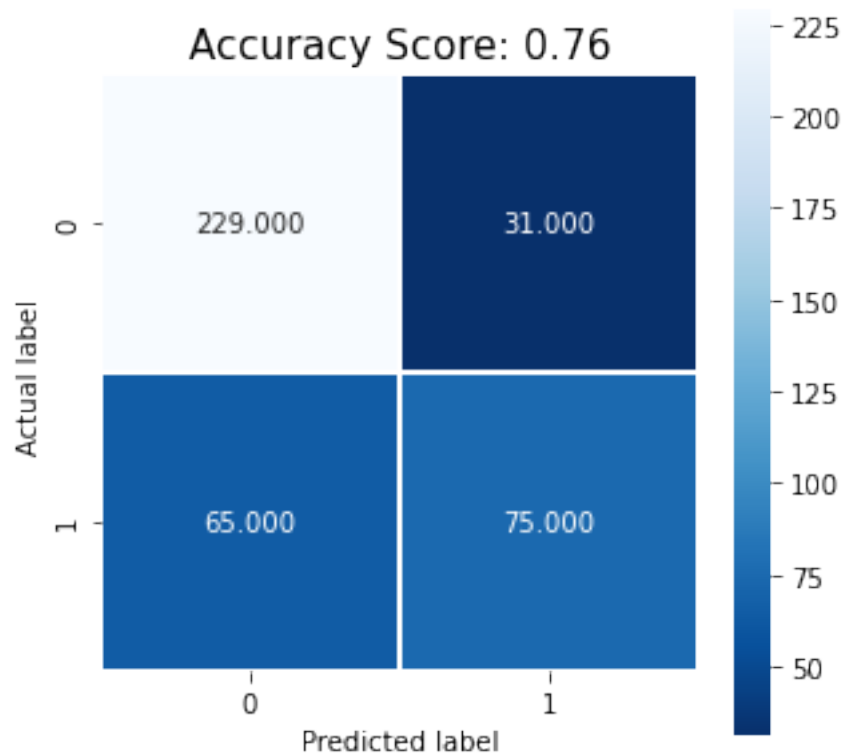
```
[17]: import matplotlib.pyplot as plt # vis
import seaborn as sns # vis / distribution

from sklearn import metrics

cm = metrics.confusion_matrix(y_test, model.predict(X_test))
```



```
[18]: plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'Blues_r');
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Accuracy Score: {0}'.format(score_test)
plt.title(all_sample_title, size = 15);
```



```
[19]: precision = cm[1,1] / (cm[1,1] + cm[0,1])
recall = cm[1,1] / (cm[1,1] + cm[1,0])
f1 = 2 * precision * recall / (precision + recall)
print('Precision: {0}, \nRecall: {1}, \nF1 Score: {2}'.
      format(precision,recall,f1))
```

Precision: 0.7075471698113207,
 Recall: 0.5357142857142857,
 F1 Score: 0.6097560975609755

Others Look at other metrics: Adjusted R^2 , Mallow's C_p , AIC, BIC etc.

1.2 The Problem of Overfitting

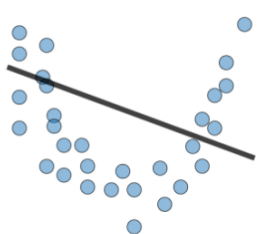
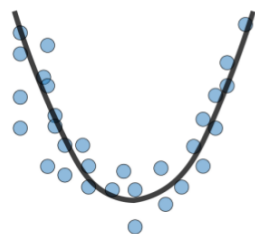
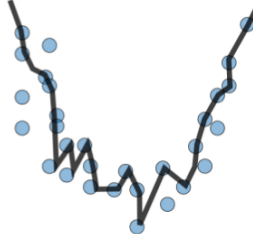
Simply put, overfitting arises when your model has fit the data too well. That can seem weird at first glance. The whole point of machine learning is to fit the data. How can it be that your model is too good at that? In machine learning, there are two really important measures you should be paying attention to at all times: the training error and the test error. We introduced them at the beginning of the notebook. Training error is a measure of how well your model performed in training, and test error is how well it performed in the wild.

When the model performs well in training data but it does not perform well in test data. We call it **overfitting**.

If we have too many features, the learned model may fit the training data set very well, but fail to generalize to new examples (test set). That means, if our model learns training set very well, we can't generalize these model for other subsets.

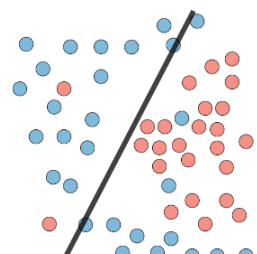
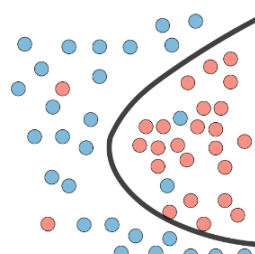
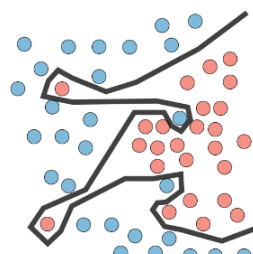
```
[8]: from IPython.display import Image
from IPython.core.display import HTML
Image(filename= "./img/of4.png",width=700, height=500)
```

[8]:

	Underfitting	Just right	Overfitting
Symptoms	<ul style="list-style-type: none"> - High training error - Training error close to test error - High bias 	<ul style="list-style-type: none"> - Training error slightly lower than test error 	<ul style="list-style-type: none"> - Low training error - Training error much lower than test error - High variance
Regression			

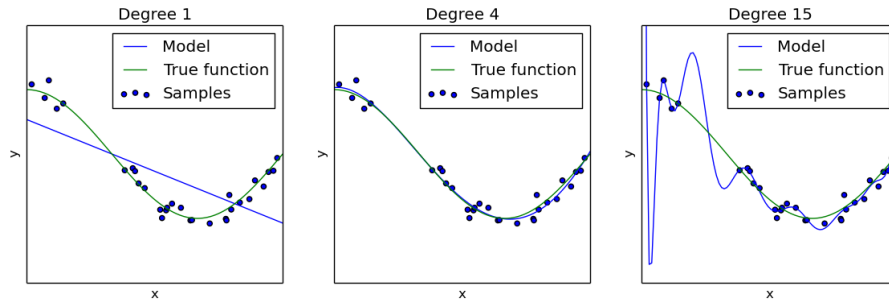
```
[9]: Image(filename= "./img/of5.png",width=700, height=500)
```

[9]:

Classification			
----------------	---	--	---

```
[14]: Image(filename= "./img/of6.png",width=1000, height=500)
```

[14]:



Bias - Variance Trade-Off Let's visualize overfitting on linear regression.

High variance means overfitting, high bias means underfitting.

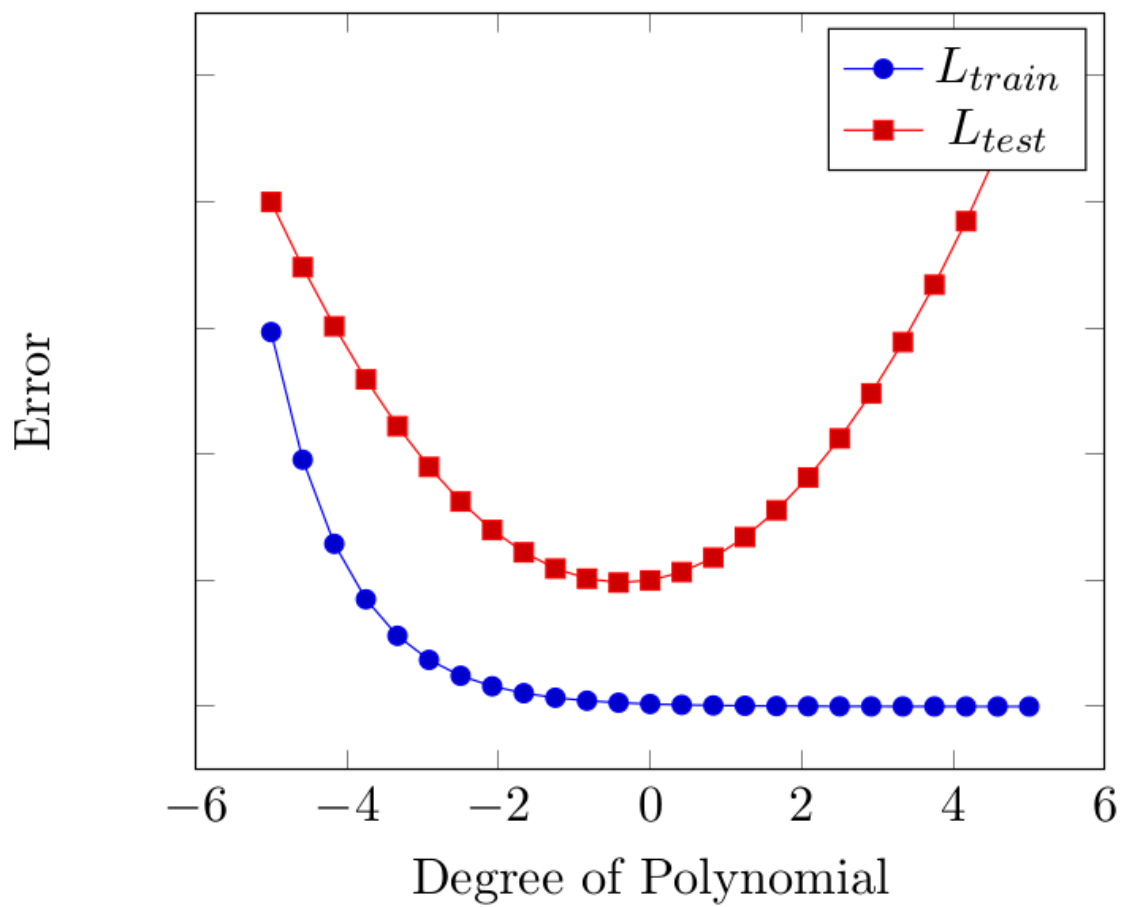
Defining training and test error:

$$L_{train} = \frac{1}{2m_{train}} \sum_i (y_{train}^i - \hat{y}_{train}^i)^2$$

$$L_{test} = \frac{1}{2m_{test}} \sum_i (y_{test}^i - \hat{y}_{test}^i)^2$$

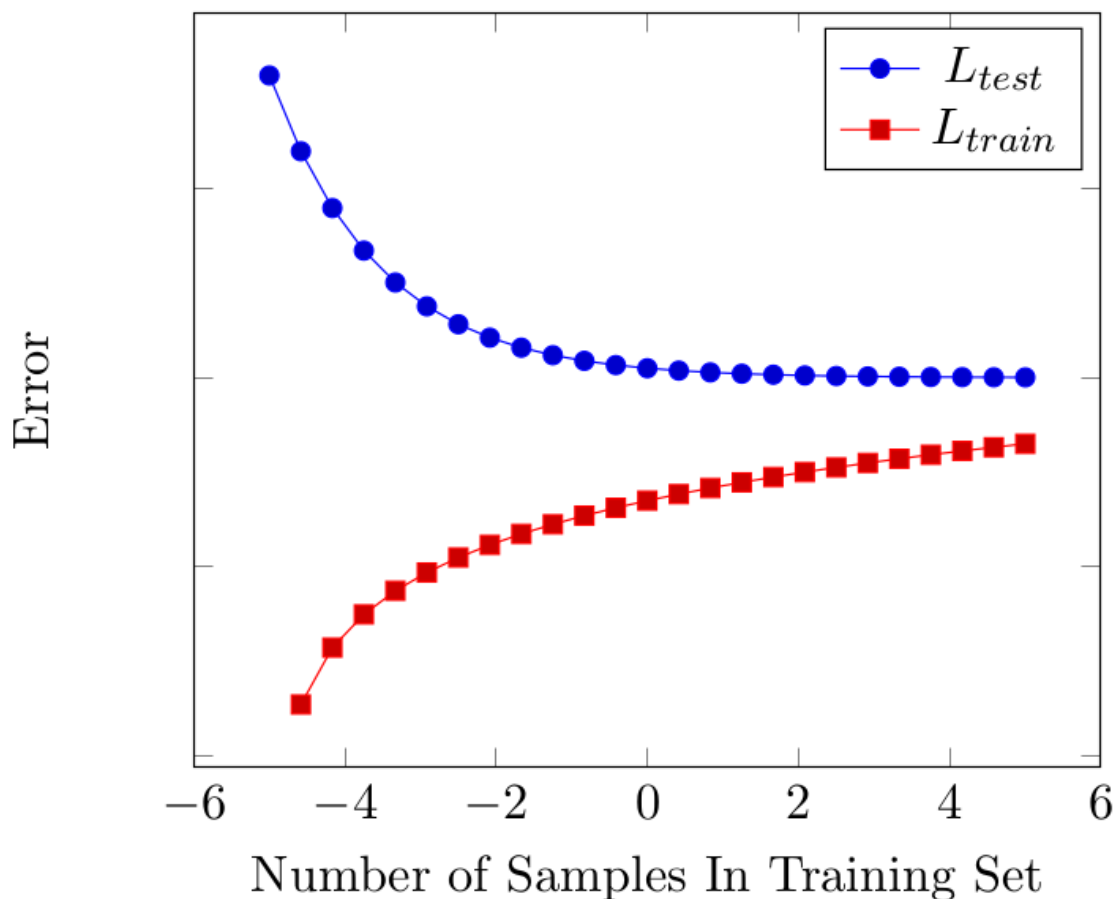
```
[19]: from IPython.display import Image
from IPython.core.display import HTML
Image(filename= "./img/biasvar1.png",width=400, height=40)
```

[19]:



```
[1]: from IPython.display import Image
      from IPython.core.display import HTML
      Image(filename= "./img/biasvar2.png",width=400, height=40)
```

[1]:



So, what we should do when overfitting happens?

- Get more training examples \rightarrow fixes high variance.
- Try smaller sets of features (look at Principal Component Analysis) \rightarrow fixes high variance.
- Try getting additional features \rightarrow fixes high bias.
- Try adding polynomial features \rightarrow fixes high bias.

So let's prove that the expected MSE, for a given observation x , can be always be decomposed into the sum of three fundamental quantities: the variance, the squared bias, the variance error ϵ .

Our objective is to, for a fixed point x , evaluate how closely the estimator can estimate the noisy observation Y corresponding to x .

Again, we can view D as the training data, and (x, y) as a test point — the test point x is probably not even in the training set D ! Mathematically, we express our metric as the expected squared error between the estimator and the observation:

$$\mathbb{E}[(f(x; D) - y)^2] = \frac{1}{m} \sum_i (y^i - \hat{y}^i)^2$$

The error metric is difficult to interpret and work with, so let's try to decompose it into parts that are easier to understand.

```
[5]: from IPython.display import Image
from IPython.core.display import HTML
Image(filename= "./img/biasvar_dec.png",width=800, height=800)
```

[5]: derivation. At its core, it uses the technique that $\mathbb{E}[(Z - Y)^2] = \mathbb{E}[(Z - \mathbb{E}[Z]) + (\mathbb{E}[Z] - Y)]^2$ which decomposes to easily give us the variance of Z and other terms.

$$\begin{aligned}
\varepsilon(\mathbf{x}; h) &= \mathbb{E}[(h(\mathbf{x}; \mathcal{D}) - Y)^2] \\
&= \mathbb{E}[(h(\mathbf{x}; \mathcal{D}) - \mathbb{E}[h(\mathbf{x}; \mathcal{D})] + \mathbb{E}[h(\mathbf{x}; \mathcal{D})] - Y)^2] \\
&= \mathbb{E}[(h(\mathbf{x}; \mathcal{D}) - \mathbb{E}[h(\mathbf{x}; \mathcal{D})])^2] + \mathbb{E}[(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - Y)^2] + 2\mathbb{E}[(h(\mathbf{x}; \mathcal{D}) - \mathbb{E}[h(\mathbf{x}; \mathcal{D})]) \cdot (\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - Y)] \\
&= \mathbb{E}[(h(\mathbf{x}; \mathcal{D}) - \mathbb{E}[h(\mathbf{x}; \mathcal{D})])^2] + \mathbb{E}[(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - Y)^2] + 2\mathbb{E}[h(\mathbf{x}; \mathcal{D}) - \mathbb{E}[h(\mathbf{x}; \mathcal{D})]] \cdot \mathbb{E}[\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - Y] \\
&= \mathbb{E}[(h(\mathbf{x}; \mathcal{D}) - \mathbb{E}[h(\mathbf{x}; \mathcal{D})])^2] + \mathbb{E}[(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - Y)^2] \\
&= \text{Var}(h(\mathbf{x}; \mathcal{D})) + \mathbb{E}[(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - Y)^2] \\
&= \text{Var}(h(\mathbf{x}; \mathcal{D})) + \mathbb{E}[(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - \mathbb{E}[Y] + \mathbb{E}[Y] - Y)^2] \\
&= \text{Var}(h(\mathbf{x}; \mathcal{D})) + \mathbb{E}[(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - \mathbb{E}[Y])^2] + \mathbb{E}[(Y - \mathbb{E}[Y])^2] + 2(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - \mathbb{E}[Y]) \cdot \mathbb{E}[\mathbb{E}[Y] - Y] \\
&= \text{Var}(h(\mathbf{x}; \mathcal{D})) + \mathbb{E}[(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - \mathbb{E}[Y])^2] + \mathbb{E}[(Y - \mathbb{E}[Y])^2] \\
&= \text{Var}(h(\mathbf{x}; \mathcal{D})) + (\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - \mathbb{E}[Y])^2 + \text{Var}(Y) \\
&= \text{Var}(h(\mathbf{x}; \mathcal{D})) + (\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - f(\mathbf{x}))^2 + \text{Var}(Z) \\
&= \underbrace{(\mathbb{E}[h(\mathbf{x}; \mathcal{D})] - f(\mathbf{x}))^2}_{\text{bias}^2 \text{ of method}} + \underbrace{\text{Var}(h(\mathbf{x}; \mathcal{D}))}_{\text{variance of method}} + \underbrace{\text{Var}(Z)}_{\text{irreducible error}}
\end{aligned}$$

[]: